

**TU Berlin**  
**Machine Learning 1**

---

**05 Homework**  
**Model Selection**

---

Jiayun, Fabi, Vincent, Lukas, Paul  
December 5, 2020

**Exercise 1: Bias and Variance of Mean Estimators****Bias:**

$$\begin{aligned} \text{Bias}(\hat{\mu}) &= E[\hat{\mu} - \mu] = E\left[\alpha \cdot \frac{1}{N} \sum_{i=1}^N x_i - \mu\right] = \alpha E\left[\frac{1}{N} \sum_{i=1}^N x_i\right] - E[\mu] = \frac{\alpha}{N} \sum_{i=1}^N \underbrace{E[x_i]}_{=\mu} - \mu = \\ &= \frac{\alpha}{N} \cdot N \cdot \mu - \mu = \alpha \cdot \mu - \mu = \mu(\alpha - 1) \end{aligned}$$

**Variance:**

$$\text{Var}(\hat{\mu}) = \text{Var}\left(\frac{\alpha}{N} \sum_{i=1}^N x_i\right) = \frac{\alpha^2}{N^2} \text{Var}\left(\sum_{i=1}^N x_i\right) = \frac{\alpha^2}{N^2} \sum_{i=1}^N \underbrace{\text{Var}(x_i)}_{=\sigma^2} = \frac{\alpha^2}{N^2} N \sigma^2 = \frac{\alpha^2}{N} \sigma^2$$

**Mean Squared Error:**

$$\text{Error}(\hat{\mu}) = \text{Bias}(\hat{\mu})^2 + \text{Var}(\hat{\mu}) = \mu^2(\alpha - 1)^2 + \frac{1}{N} \alpha^2 \sigma^2$$

## Exercise 2

a) First we reparameterize  $R_i = \exp(Z_i)$  to ensure that  $R$  is positive. To find the distribution that minimizes its expected KL divergence from the the class distribution estimator we have to set the derivative to 0. Observe that the derivative of  $Z_j$  is

$$\begin{aligned} \frac{d}{dZ_j} E[D_{KL}(R||\hat{P})] &= E[\frac{d}{dZ_j} D_{KL}(R||\hat{P})] = E[\frac{d}{dZ_j} \sum_{i=1}^C \exp(Z_i) \log(\exp(Z_i)/\hat{P}_i)] \\ &= E[\frac{d}{dZ_j} \sum_{i=1}^C \exp(Z_i)(Z_i - \log(\hat{P}_i))] = E[\exp(Z_j)(Z_j - \log(\hat{P}_j)) + \exp(Z_j)] \\ &= \exp(Z_j)(Z_j - E[\log(\hat{P}_j)] + 1) \end{aligned}$$

for all  $1 \leq j \leq C$ . Now setting this equal to zero we get

$$Z_j = E[\log(\hat{P}_j)] + 1 \implies R_j = \exp(E[\log(\hat{P}_j)] + 1)$$

To get the actual probability distribution we have to normalize our solution. Thus

$$R_i = \frac{\exp(E[\log(\hat{P}_i)]) \cdot e}{\sum_{j=1}^C \exp(E[\log(\hat{P}_j)]) \cdot e} = \frac{\exp(E[\log(\hat{P}_i)])}{\sum_{j=1}^C \exp(E[\log(\hat{P}_j)])}$$

for  $1 \leq i \leq C$ .

b) First we consider the term that is given in the hint, which will be useful in our calculations

$$\begin{aligned} E[\log(R_i) - \log(\hat{P}_i)] &= E[E[\log(\hat{P}_i)] - \log(\sum_{j=1}^C \exp(E[\log(\hat{P}_j)])) - \log(\hat{P}_i)] \\ &= -\log(\sum_{j=1}^C \exp(E[\log(\hat{P}_j)])) =: E[\log(R) - \log(\hat{P})] \end{aligned}$$

So this does indeed not depend on  $i$ . One can show with the same calculations that this also works for

$$E[\log(P_i) - \log(R_i)] =: E[\log(P) - \log(R)]$$

We can write by definition of the expectation with respect to the probability distribution defined over the simplex  $[1, \dots, C]$  as

$$\mathbb{E}[\log(P) - \log(R)] = \sum_{i=1}^C P_i(\log(P_i) - \log(R_i))$$

and

$$\mathbb{E}[\log(R) - \log(\hat{P})] = \sum_{i=1}^C R_i(\log(R_i) - \log(\hat{P}_i))$$

Now we can show the bias-variance decomposition

$$\begin{aligned} \text{Bias}(\hat{P}) + \text{Var}(\hat{P}) &= D_{KL}(P||R) + \mathbb{E}[D_{KL}(R||\hat{P})] \\ &= \sum_{i=1}^C P_i \log(P_i/R_i) + \mathbb{E}[\sum_{i=1}^C R_i \log(R_i/\hat{P}_i)] \\ &= \sum_{i=1}^C P_i(\log(P_i) - \log(R_i)) + \mathbb{E}[\sum_{i=1}^C R_i(\log(R_i) - \log(\hat{P}_i))] \\ &= \mathbb{E}[\log(P) - \log(R)] + \mathbb{E}[\mathbb{E}[\log(R) - \log(\hat{P})]] \\ &= \mathbb{E}[\mathbb{E}[\log(P) - \log(R)]] + \mathbb{E}[\log(R) - \log(\hat{P})] \\ &= \mathbb{E}[\mathbb{E}[\log(P) - \log(\hat{P})]] \\ &= \mathbb{E}[\sum_{i=1}^C P_i(\log(P_i) - \log(\hat{P}_i))] \\ &= \mathbb{E}[D_{KL}(P||\hat{P})] = \text{Error}(\hat{P}) \end{aligned}$$

# sheet05

December 6, 2020

## 0.1 Part 1: The James-Stein Estimator (20 P)

Let  $x_1, \dots, x_N \in \mathbb{R}^d$  be independent draws from a multivariate Gaussian distribution with mean vector  $\mu$  and covariance matrix  $\Sigma = \sigma^2 I$ . It can be shown that the maximum-likelihood estimator of the mean parameter  $\mu$  is the empirical mean given by:

$$\hat{\mu}_{\text{ML}} = \frac{1}{N} \sum_{i=1}^N x_i$$

Maximum-likelihood appears to be a strong estimator. However, it was demonstrated that the following estimator

$$\hat{\mu}_{JS} = \left(1 - \frac{(d-2) \cdot \frac{\sigma^2}{N}}{\|\hat{\mu}_{\text{ML}}\|^2}\right) \hat{\mu}_{\text{ML}}$$

(a shrunk version of the maximum-likelihood estimator towards the origin) has actually a smaller distance from the true mean when  $d \geq 3$ . This however assumes knowledge of the variance of the distribution for which the mean is estimated. This estimator is called the James-Stein estimator. While the proof is a bit involved, this fact can be easily demonstrated empirically through simulation. This is the object of this exercise.

The code below draws ten 50-dimensional points from a normal distribution with mean vector  $\mu = (1, \dots, 1)$  and covariance  $\Sigma = I$ .

```
[123]: import numpy

def getdata(seed):

    n = 10                # data points
    d = 50                # dimensionality of data
    m = numpy.ones([d])   # true mean
    s = 1.0               # true standard deviation

    rstate = numpy.random.mtrand.RandomState(seed)
    X = rstate.normal(0,1,[n,d])*s+m

    return X,m,s
```

The following function computes the maximum likelihood estimator from a sample of the data assumed to be generated by a Gaussian distribution:

```
[124]: def ML(X):
        return X.mean(axis=0)
```

### 0.1.1 Implementing the James-Stein Estimator (10 P)

- Based on the ML estimator function, write a function that receives as input the data  $(X_i)_{i=1}^n$  and the (known) variance  $\sigma^2$  of the generating distribution, and computes the James-Stein estimator

```
[125]: def JS(X,s):
        # REPLACE BY YOUR CODE
        Uml = ML(X)
        (N,d) = X.shape
        m_JS = (1 - (d-2)*s**2/(numpy.linalg.norm(Uml,2)**2*N))*Uml
        ###
        return m_JS
```

### 0.1.2 Comparing the ML and James-Stein Estimators (10 P)

We would like to compute the error of the maximum likelihood estimator and the James-Stein estimator for 100 different samples (where each sample consists of 10 draws generated by the function `getdata` with a different random seed). Here, for reproducibility, we use seeds from 0 to 99. The error should be measured as the Euclidean distance between the true mean vector and the estimated mean vector.

- Compute the maximum-likelihood and James-Stein estimations.
- Measure the error of these estimations.
- Build a scatter plot comparing these errors for different samples.

```
[126]: %matplotlib inline
        ### REPLACE BY YOUR CODE
        from matplotlib import pyplot as plt
        Error_ml = []
        Error_js = []
        for seed in range(100):
            (X,m,s) = getdata(seed)
            m_ML = ML(X)
            m_JS = JS(X,s)


            Error_ml.append(numpy.linalg.norm((m_ML-1), 2))
            Error_js.append(numpy.linalg.norm((m_JS-1), 2))

        Error_js = numpy.array(Error_js)
        Error_ml = numpy.array(Error_ml)

        plt.figure(figsize=(6,6))
```

```
plt.scatter(Error_ml, Error_js,c = 'k')
plt.xlim(0, 5)
plt.ylim(0, 5)
plt.xlabel('Error for Maximum Likelihood estimator')
plt.ylabel('Error for James Stein estimator')
###
```

[126]: `Text(0, 0.5, 'Error for James Stein estimator')`



## 0.2 Part 2: Bias/Variance Decomposition (30 P)

In this part, we would like to implement a procedure to find the bias and variance of different predictors. We consider one for regression and one for classification. These predictors are available in the module `utils`.

- **`utils.ParzenRegressor`:** A regression method based on Parzen window. The hyperparam-

eter corresponds to the scale of the Parzen window. A large scale creates a more rigid model. A small scale creates a more flexible one.

- **utils.ParzenClassifier:** A classification method based on Parzen window. The hyperparameter corresponds to the scale of the Parzen window. A large scale creates a more rigid model. A small scale creates a more flexible one. Note that instead of returning a single class for a given data point, it outputs a probability distribution over the set of possible classes.

Each class of predictor implements the following three methods:

- **\_\_init\_\_(self,parameter):** Create an instance of the predictor with a certain scale parameter.
- **fit(self,X,T):** Fit the predictor to the data (a set of data points  $X$  and targets  $T$ ).
- **predict(self,X):** Compute the output values arbitrary inputs  $X$ .

To compute the bias and variance estimates, we require *multiple samples* from the training set for a single set of observation data. To accomplish this, we utilize the **Sampler** class provided. The sampler is initialized with the training data and passed to the method for estimating bias and variance, where its function **sampler.sample()** is called repeatedly in order to fit multiple models and create an ensemble of prediction for each test data point.

### 0.2.1 Regression Case (15 P)

For the regression case, Bias, Variance and Error are given by:

- $\text{Bias}(Y)^2 = (\mathbb{E}_Y[Y - T])^2$
- $\text{Var}(Y) = \mathbb{E}_Y[(Y - \mathbb{E}_Y[Y])^2]$
- $\text{Error}(Y) = \mathbb{E}_Y[(Y - T)^2]$

**Task:** Implement the KL-based Bias-Variance Decomposition defined above. The function should repeatedly sample training sets from the sampler (as many times as specified by the argument `nbsamples`), learn the predictor on them, and evaluate the variance on the out-of-sample distribution given by  $X$  and  $T$ .

```
[127]: def biasVarianceRegression(sampler, predictor, X, T, nbsamples):

    # -----
    # TODO: REPLACE BY YOUR CODE
    # -----
    Y_sample = []
    for _ in range(nbsamples):
        sample = sampler.sample()
        predictor.fit(sample[0],sample[1])
        Y_sample.append(predictor.predict(X))
    Y_sample = numpy.array(Y_sample)
    Y = numpy.mean(Y_sample,axis=0)

    variance = numpy.mean(numpy.mean((Y_sample - Y)**2,axis=0))
    error = numpy.mean(numpy.mean((Y_sample - T)**2,axis=0))
```



```

bias = error - variance
# -----

return bias, variance


```

Your implementation can be tested with the following code:

```

[128]: import utils, numpy
        %matplotlib inline
        utils.plotBVE(utils.Housing, numpy.logspace(-6, 3, num=30), utils.
        ↪ ParzenRegressor, biasVarianceRegression, 'Housing Regression')

```



## 0.2.2 Classification Case (15 P)

We consider here the Kullback-Leibler divergence as a measure of classification error, as derived in the exercise, the Bias, Variance decomposition for such error is:

- $\text{Bias}(Y) = D_{\text{KL}}(T||R)$
- $\text{Var}(Y) = \mathbb{E}_Y[D_{\text{KL}}(R||Y)]$
- $\text{Error}(Y) = \mathbb{E}_Y[D_{\text{KL}}(T||Y)]$

where  $R$  is the distribution that minimizes its expected KL divergence from the estimator of probability distribution  $Y$  (see the theoretical exercise for how it is computed exactly), and where  $T$  is the target class distribution.

**Task:** Implement the KL-based Bias-Variance Decomposition defined above. The function should repeatedly sample training sets from the sampler (as many times as specified by the argument `nbsamples`), learn the predictor on them, and evaluate the variance on the out-of-sample distribution given by `X` and `T`.

```
[129]: def KL(P1,P2):
        return numpy.sum(P1 * numpy.log(P1 / P2), axis=2)

def biasVarianceClassification(sampler, predictor, X, T, nbsamples=25):

    # -----
    # TODO: REPLACE BY YOUR CODE
    # -----
    Y_sample = []
    for _ in range(nbsamples):
        sample = sampler.sample()
        predictor.fit(sample[0],sample[1])
        Y_sample.append(predictor.predict(X))
    Y_sample = numpy.array(Y_sample) #(20, 742, 7)

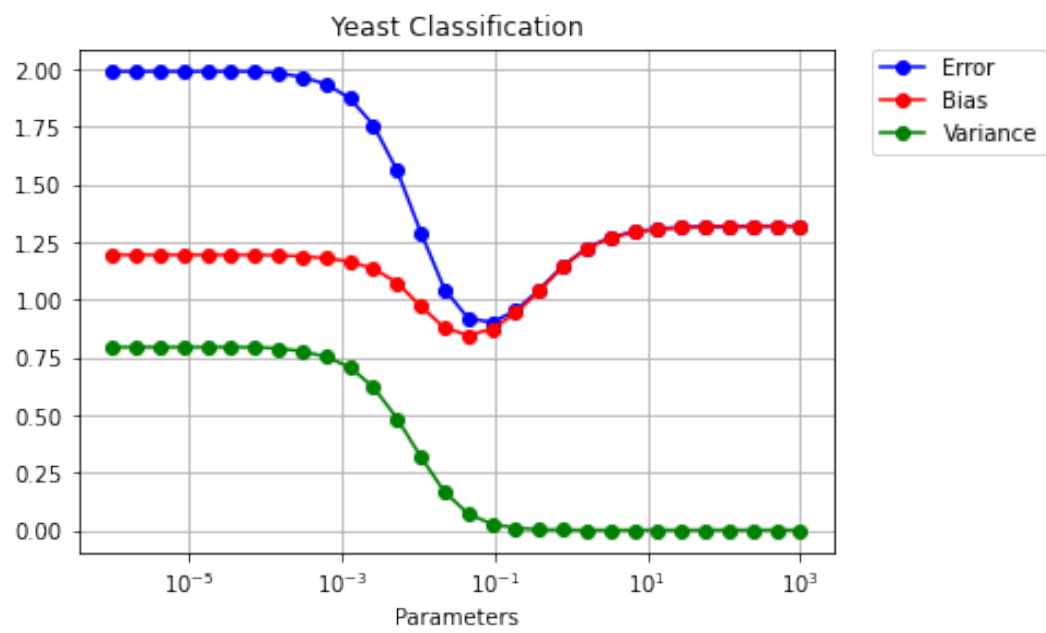
    numerator = numpy.exp(numpy.mean(numpy.log(Y_sample),axis=0)) #(742,7)
    denominator = numpy.sum(numerator,axis=1)[: ,numpy.newaxis] #(742,1)
    R = numerator / denominator #(742,7)

    Error = numpy.mean(KL(T,Y_sample)) #scalar
    variance = numpy.mean(KL(R,Y_sample)) #inner KL is (20,742),variance is scalar
    bias = Error - variance
    # -----

    return bias,variance
```

Your implementation can be tested with the following code:

```
[130]: import utils,numpy
        %matplotlib inline
        utils.plotBVE(utils.Yeast,numpy.logspace(-6,3,num=30),utils.
        ↪ParzenClassifier,biasVarianceClassification,'Yeast Classification')
```



[ ]: